

problem statement:-To Predict the Rainfall based on the various features of the dataset

import the Essential Libraries

```
In [2]: import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: df=pd.read_csv(r"C:\Users\Venky\AppData\Local\Temp\Temp1_100Years_RainfallDataset.zip\rainfall in india 1901-2015.csv")
df
```

Out[3]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	Mar-May	Jun-Sep	Oct-Dec
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5	558.2	33.6	3373.2	136.3	560.3	1696.3	980
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2	359.0	160.5	3520.7	159.8	458.3	2185.9	716
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2	284.4	225.0	2957.4	156.7	236.1	1874.0	690
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2	308.7	40.1	3079.6	24.1	506.9	1977.6	571
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7	25.4	344.7	2566.7	1.3	309.7	1624.9	630
...
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4	184.3	14.9	1533.7	7.9	196.2	1013.0	316
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9	12.4	8.8	1405.5	19.3	99.6	1119.5	167
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8	78.1	26.7	1426.3	60.6	131.1	1057.0	177
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.2	59.0	62.3	1395.0	69.3	76.7	958.5	290
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4	231.0	159.0	1642.9	2.7	223.9	860.9	555

4116 rows × 19 columns

Data preprocessing

In [4]: `df.head()`

Out[4]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	Mar-May	Jun-Sep	Oct-Dec
0	ANDAMAN & NICOBAR ISLANDS	1901	49.2	87.1	29.2	2.3	528.8	517.5	365.1	481.1	332.6	388.5	558.2	33.6	3373.2	136.3	560.3	1696.3	980.3
1	ANDAMAN & NICOBAR ISLANDS	1902	0.0	159.8	12.2	0.0	446.1	537.1	228.9	753.7	666.2	197.2	359.0	160.5	3520.7	159.8	458.3	2185.9	716.7
2	ANDAMAN & NICOBAR ISLANDS	1903	12.7	144.0	0.0	1.0	235.1	479.9	728.4	326.7	339.0	181.2	284.4	225.0	2957.4	156.7	236.1	1874.0	690.6
3	ANDAMAN & NICOBAR ISLANDS	1904	9.4	14.7	0.0	202.4	304.5	495.1	502.0	160.1	820.4	222.2	308.7	40.1	3079.6	24.1	506.9	1977.6	571.0
4	ANDAMAN & NICOBAR ISLANDS	1905	1.3	0.0	3.3	26.9	279.5	628.7	368.7	330.5	297.0	260.7	25.4	344.7	2566.7	1.3	309.7	1624.9	630.8

In [5]: `df.tail()`

Out[5]:

	SUBDIVISION	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	OCT	NOV	DEC	ANNUAL	Jan-Feb	Mar-May	Jun-Sep	Oct-Dec
4111	LAKSHADWEEP	2011	5.1	2.8	3.1	85.9	107.2	153.6	350.2	254.0	255.2	117.4	184.3	14.9	1533.7	7.9	196.2	1013.0	316.6
4112	LAKSHADWEEP	2012	19.2	0.1	1.6	76.8	21.2	327.0	231.5	381.2	179.8	145.9	12.4	8.8	1405.5	19.3	99.6	1119.5	167.1
4113	LAKSHADWEEP	2013	26.2	34.4	37.5	5.3	88.3	426.2	296.4	154.4	180.0	72.8	78.1	26.7	1426.3	60.6	131.1	1057.0	177.6
4114	LAKSHADWEEP	2014	53.2	16.1	4.4	14.9	57.4	244.1	116.1	466.1	132.2	169.2	59.0	62.3	1395.0	69.3	76.7	958.5	290.5
4115	LAKSHADWEEP	2015	2.2	0.5	3.7	87.1	133.1	296.6	257.5	146.4	160.4	165.4	231.0	159.0	1642.9	2.7	223.9	860.9	555.4

```
In [6]: df.isnull().any()
```

```
Out[6]: SUBDIVISION    False  
YEAR                False  
JAN                  True  
FEB                  True  
MAR                  True  
APR                  True  
MAY                  True  
JUN                  True  
JUL                  True  
AUG                  True  
SEP                  True  
OCT                  True  
NOV                  True  
DEC                  True  
ANNUAL               True  
Jan-Feb              True  
Mar-May              True  
Jun-Sep              True  
Oct-Dec              True  
dtype: bool
```

```
In [7]: df.fillna(method='ffill',inplace=True)
```

```
In [8]: df.fillna(method='bfill',inplace=True)
```

```
In [9]: df.isnull().sum()
```

```
Out[9]: SUBDIVISION    0  
        YEAR          0  
        JAN           0  
        FEB           0  
        MAR           0  
        APR           0  
        MAY           0  
        JUN           0  
        JUL           0  
        AUG           0  
        SEP           0  
        OCT           0  
        NOV           0  
        DEC           0  
        ANNUAL        0  
        Jan-Feb       0  
        Mar-May       0  
        Jun-Sep       0  
        Oct-Dec       0  
        dtype: int64
```

In [10]: df.describe()

Out[10]:

	YEAR	JAN	FEB	MAR	APR	MAY	JUN	JUL	AUG	SEP	
count	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000	4116.000000
mean	1958.218659	18.957240	21.823251	27.415379	43.160641	85.788994	230.567979	347.177235	290.239796	197.524781	95.724781
std	33.140898	33.576192	35.922602	47.045473	67.816588	123.220150	234.896056	269.321089	188.785639	135.509037	99.609037
min	1901.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.400000	0.000000	0.000000	0.100000	0.000000
25%	1930.000000	0.600000	0.600000	1.000000	3.000000	8.600000	70.475000	175.900000	155.850000	100.575000	14.600000
50%	1958.000000	6.000000	6.700000	7.900000	15.700000	36.700000	138.900000	284.800000	259.400000	174.000000	65.724781
75%	1987.000000	22.200000	26.800000	31.400000	50.125000	97.400000	306.150000	418.325000	377.800000	266.225000	148.609037
max	2015.000000	583.700000	403.500000	605.600000	595.100000	1168.600000	1609.900000	2362.800000	1664.600000	1222.000000	948.300000

```
In [11]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4116 entries, 0 to 4115
Data columns (total 19 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   SUBDIVISION     4116 non-null  object 
 1   YEAR            4116 non-null  int64  
 2   JAN             4116 non-null  float64
 3   FEB             4116 non-null  float64
 4   MAR             4116 non-null  float64
 5   APR             4116 non-null  float64
 6   MAY             4116 non-null  float64
 7   JUN             4116 non-null  float64
 8   JUL             4116 non-null  float64
 9   AUG             4116 non-null  float64
10  SEP             4116 non-null  float64
11  OCT             4116 non-null  float64
12  NOV             4116 non-null  float64
13  DEC             4116 non-null  float64
14  ANNUAL          4116 non-null  float64
15  Jan-Feb         4116 non-null  float64
16  Mar-May         4116 non-null  float64
17  Jun-Sep         4116 non-null  float64
18  Oct-Dec         4116 non-null  float64
dtypes: float64(17), int64(1), object(1)
memory usage: 611.1+ KB
```

```
In [12]: df.columns
```

```
Out[12]: Index(['SUBDIVISION', 'YEAR', 'JAN', 'FEB', 'MAR', 'APR', 'MAY', 'JUN', 'JUL',
               'AUG', 'SEP', 'OCT', 'NOV', 'DEC', 'ANNUAL', 'Jan-Feb', 'Mar-May',
               'Jun-Sep', 'Oct-Dec'],
              dtype='object')
```

```
In [13]: df.shape
```

```
Out[13]: (4116, 19)
```

```
In [15]: df['ANNUAL'].value_counts()
```

```
Out[15]: ANNUAL
790.5      4
770.3      4
1836.2     4
1024.6     4
1926.5     3
..
443.9      1
689.0      1
605.2      1
509.7      1
1642.9     1
Name: count, Length: 3712, dtype: int64
```

```
In [17]: df['Jan-Feb'].value_counts()
```

```
Out[17]: Jan-Feb
0.0      238
0.1       80
0.2       52
0.3       38
0.4       32
...
23.3       1
95.2       1
76.9       1
66.5       1
69.3       1
Name: count, Length: 1220, dtype: int64
```



```
In [18]: df['Mar-May'].value_counts()
```

```
Out[18]: Mar-May
0.0      29
0.1      13
0.3      11
8.3       11
11.5      10
..
246.3     1
248.1     1
151.3     1
249.5     1
223.9     1
Name: count, Length: 2262, dtype: int64
```

```
In [19]: df['Jun-Sep'].value_counts()
```


```
Out[19]: Jun-Sep
434.3     4
334.8     4
573.8     4
613.3     4
1082.3    3
..
301.6     1
380.9     1
409.3     1
229.4     1
958.5     1
Name: count, Length: 3683, dtype: int64
```

```
In [20]: df['Oct-Dec'].value_counts()
```

```
Out[20]: Oct-Dec  
0.0      16  
0.1      15  
0.5      13  
0.6      12  
0.7      11  
..  
191.5     1  
124.5     1  
139.1     1  
41.5      1  
555.4     1  
Name: count, Length: 2389, dtype: int64
```

Exploratory Data analysis

```
In [21]: df=df[['JAN', 'FEB', 'MAR', 'APR', 'DEC']]  
sns.heatmap(df.corr(),annot=True)  
plt.show()
```



```
In [22]: df.columns
```

```
Out[22]: Index(['JAN', 'FEB', 'MAR', 'APR', 'DEC'], dtype='object')
```

```
In [23]: x=df[['FEB']]  
y=df[['JAN']]
```

Linear Regression

```
In [24]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [25]: from sklearn.linear_model import LinearRegression
reg=LinearRegression()
reg.fit(x_train,y_train)
print(reg.intercept_)
coeff_=pd.DataFrame(reg.coef_,x.columns,columns=['coefficient'])
coeff_
```

[9.65066661]

Out[25]:

	<u>coefficient</u>
FEB	0.442278

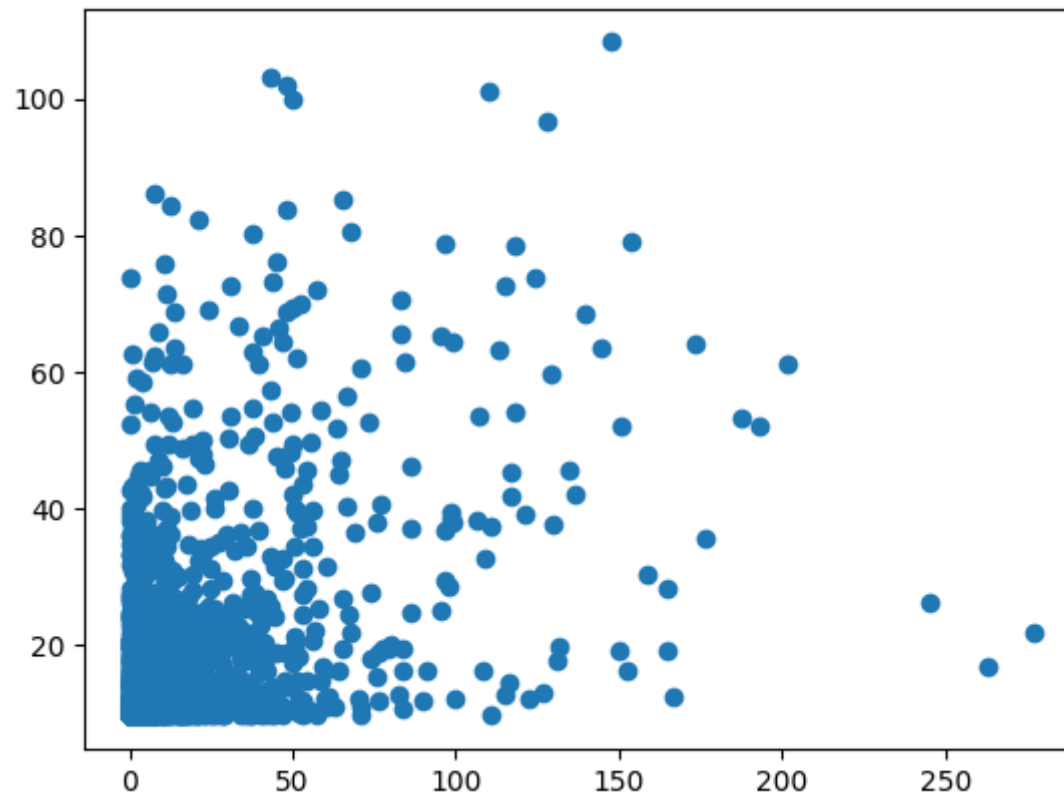
```
In [26]: score=reg.score(x_test,y_test)
print(score)
```

0.1793580786264921


```
In [28]: predictions=reg.predict(x_test)
```

```
In [29]: plt.scatter(y_test,predictions)
```

```
Out[29]: <matplotlib.collections.PathCollection at 0x1e3fb64e9d0>
```




```
In [30]: df500=df[:][:500]  
sns.lmplot(x='FEB',y='JAN',order=2,ci=None,data=df500)  
plt.show()
```



```
In [32]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33)
reg.fit(x_train,y_train)
reg.fit(x_test,y_test)
```

```
Out[32]: ▼ LinearRegression
LinearRegression()
```

```
In [34]: y_pred=reg.predict(x_test)
plt.scatter(x_test,y_test,color='green')
plt.plot(x_test,y_pred,color='red')
plt.show()
```



```
In [36]: from sklearn.linear_model import Lasso,Ridge  
from sklearn.preprocessing import StandardScaler
```

```
In [37]: features=df.columns[0:5]  
targets=df.columns[-5]
```

```
In [38]: x=df[features].values  
y=df[targets].values  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=17)
```

```
In [39]: ridgeReg=Ridge(alpha=10)  
ridgeReg.fit(x_train,y_train)  
train_score_ridge=ridgeReg.score(x_train,y_train)  
test_score_ridge=ridgeReg.score(x_test,y_test)
```

```
In [40]: print("\n Ridge Model:\n")  
print(" the train score for Ridge model is{}".format(train_score_ridge))  
print(" the test score for Ridge model is{}".format(test_score_ridge))
```


Ridge Model:

the train score for Ridge model is0.999999999874192
the test score for Ridge model is0.9999999998833

```
In [42]: lr=LinearRegression()
```



```
In [45]: plt.figure(figsize=(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize='5',color='red',label=r'Ridge,$\alpha$')
plt.plot(features,alpha=0.4,linestyle='none',marker='o',markersize='7',color='green',label=r'LinearRegression')
plt.xticks(rotation=90)
plt.legend()
plt.show()
```



Lasso Model


```
In [47]: print("\n lasso Model:\n")
lasso=Lasso(alpha=10)
lasso.fit(x_train,y_train)
train_score_ls=lasso.score(x_train,y_train)
test_score_ls=lasso.score(x_test,y_test)
print(" the train score for ls model is{}".format(train_score_ls))
print(" the test score for ls model is{}".format(test_score_ls))
```

lasso Model:

the train score for ls model is0.9999207747038827
the test score for ls model is0.9999206791315256

```
In [49]: pd.Series(lasso.coef_,features).sort_values(ascending=True).plot(kind="bar")
```

```
Out[49]: <Axes: >
```





```
In [51]: from sklearn.linear_model import LassoCV
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(x_train,y_train)
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```

```
0.9999999999999921
```

```
0.9999999999999921
```

```
In [53]: plt.figure(figsize=(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge,$\alpha$')
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'lasso;$\alpha$=grid$')
plt.plot(features,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label=r'LinearRegression')
plt.xticks(rotation=90)
plt.legend()
plt.show()
```



Elastic Net:-

```
In [54]: from sklearn.linear_model import ElasticNet
elnet=ElasticNet()
elnet.fit(x,y)
print(elnet.coef_)
print(elnet.intercept_)
print(elnet.score(x,y))
```

```
[9.99098574e-01 0.00000000e+00 3.02728910e-05 0.00000000e+00
 0.00000000e+00]
0.01625860696662329
0.9999992160905338
```

```
In [57]: y_pred_elastic=elnet.predict(x_train)
mean_squared_score=np.mean((y_pred_elastic - y_train)**2)
print(mean_squared_score)
```

```
0.0008816302333966198
```

Conclusion

among all models lasso yield highest accuracy.so we prefer Lasso model for this Rain fall dataset

In []: